

# DynamoDB Design Patterns

---



**Ivan Mushketyk**

@mushketyk   brewing.codes



# Overview



DynamoDB best practices


Data Modeling with DynamoDB

Hot Keys

Reducing cost with DynamoDB



# Data Modeling in DynamoDB



## Data modeling for

- 1:1 relationships
- 1:N relationships
- M:N relationships

## Composite keys

# 1:1 Relationship

OrderId	ItemId	Date	TotalPrice
1	1	2017_05_11	117.8
2	2	2017_08_11	167
3	2	2017_11_22	167

OrderId	Status	Courier
1	DELIVERED	FedEx
2	IN_TRANSIT	DHL
3	DELIVERED	DHL



# Why Use 1:1 Relationship



**DynamoDB item limitation**

**Update an attribute less costly**

**Can create more indexes**

**Can help to save money**

# 1:M Relationship

Partitionkey

OrderId	ItemId	Name	Price
1	4	DynamoDB sticker	1
1	5	DynamoDB book	25
2	9	Quadrocopter	300
3	4	Phone	10
3	5	Bitcoin miner	454

Sort key



# M:N Relationship

Partition  
key

Sort key

**OrderId**

**ItemId**

**Name**

**Price**

1	4	DynamoDB sticker	1
1	5	DynamoDB book	25
2	9	Quadrocopter	300
3	4	Phone	10
3	5	Bitcoin miner	454

GSI sort  
key

GSI partition key



# Composite Keys



**Can query using only one index**

**Find all orders delivered in May**

**Need to search by status AND date**

**Can use composite keys**



# Composite Keys Example

OrderId	ItemId	Status	Date
1	1	DELIVERED	2017_01_02
2	1	DELIVERED	2017_05_02
3	2	SHIPPED	2017_05_08
4	2	CANCELED	2017_06_11
5	1	DELIVERED	2017_05_14



# Composite Keys Example

OrderId	ItemId	Status	Date	Status_Date
1	1	DELIVERED	2017_01_02	DELIVERED_2017_01_02
2	1	DELIVERED	2017_05_02	DELIVERED_2017_05_02
3	2	SHIPPED	2017_05_08	SHIPPED_2017_05_08
4	2	CANCELED	2017_06_11	CANCELED_2017_06_11
5	1	DELIVERED	2017_05_14	DELIVERED_2017_05_14



# Composite Keys Example

OrderId	ItemId	Status	Date	Status_Date
1	1	DELIVERED	2017_01_02	DELIVERED_2017_01_02
2	1	DELIVERED	2017_05_02	DELIVERED_2017_05_02
3	2	SHIPPED	2017_05_08	SHIPPED_2017_05_08
4	2	CANCELED	2017_06_11	CANCELED_2017_06_11
5	1	DELIVERED	2017_05_14	DELIVERED_2017_05_14

GSI Partition key

Sort key

## Query:

Status\_Date **BEGINS\_WITH** "DELIVERED\_2017\_05"



# Hot Keys



**What are hot keys**

**Symptoms of hot keys**

**How to deal with them**

**Selecting good partition keys**

# Will This Work?




**Your table has 1000 RCUs**

**It receives 500 RCUs of requests**


**Any reasons to worry?**

# Hot Keys


Total RCU: 1000



# Avoid Hot Keys



Ensure uniform load on all partitions




Use caching




Select good partition key



# How to Select Partition Key



Boolean value



Limited range of values



UUIDs




Numbers (range is unlimited)





# Randomized Values



key="AMZN" + rand(10)



# Reducing Costs of DynamoDB



High RCU and WCU lead to high cost

How to reduce RCUs/WCUs



# Use Less Data



**Store big items in S3**

**Use data compression**

**Attributes projection**

**Split into big items into multiple tables**

# Minimize RCUs

PostId	Time	Text
1	1498916052	Let me tell the story of my life...
2	1498185048	Once upon a time...



PostId	Time	Excerpt
1	1498916052	Let me tell
2	1498185048	Once upon

PostId	Text
1	Let me tell the story of my life...
2	Once upon a time...

# Exploit Temporal Access Patterns

Tables

2018\_April

RCU: 2000  
WCU:2000

2018\_March

RCU: 1000  
WCU:1

2018\_February

RCU: 10  
WCU:1



# Time-to-live



**DynamoDB removes item when they expire**

**No extra cost**

**Allows to save money**


**No need to manually remove items**



# Time-to-live Use Cases



Legal requirements



Session data



Temporary data



# Enable Time-to-Live

```
TimeToLiveSpecification ttls = new TimeToLiveSpecification()  
ttls.setAttributeName("TTL"); ttls.setEnabled(true);
```

```
UpdateTimeToLiveRequest request = new UpdateTimeToLiveRequest();  
request.setTableName("Orders");  
request.setTimeToLiveSpecification(ttls);
```

```
client.updateTimeToLive(ttls);
```





# Sparse Indexes

## Table

OrderId    ItemId    OrderType

1	3	
2	4	
3	5	Collection
4	6	

GSI Sort key

GSI Partition key

## Index

OrderType    OrderId    ItemId

Collection	3	5
------------	---	---



# Other Tactics



Avoid scans (expensive)

Use queries

Distribute read/write operations

Reserved capacity



# Summary



NoSQL database should be used differently

How to use indexes to represent relationships

How to avoid hot keys

How to save money using DynamoDB

